

KŪPUNA-AI Bench "Let's talk story.": Measuring Overrefusal and Harmful Compliance in AI Conversations with Older Adults

Program: Anthropic Wellbeing Research Grants (\$5M pool; applications due Sep 21, 2026; full-proposal notifications Oct 5) · **Applicant:** The Gerontechnology Group, LLC (Houston, Texas) · **PI:** Melissa Mansfield, PhD, NAPG-CPG · **Technical Lead and Co-Applicant:** Jeremy A. Gracey, MS
Brief v3.3 · 2026-09-09 (applicant and team)

Items marked **[proposed]** are defaults for discussion, not decisions.

Why the name. *Kūpuna* is the Hawaiian word for grandparents, elders, and ancestors. It is a role, not an age bracket: *kūpuna* are the keepers of knowledge whose judgment a family seeks out rather than second-guesses. That is the posture this eval demands of a model: treat the older adult as the authority on their own life, and measure every refusal or hedge against that standard. The name is also where the team comes from. Both the PI and the technical lead spent years living and working in Hawai‘i; the technical lead's family has roots and a home there; and the team's medical director still practices there. Scenario variants will include Hawai‘i *kūpuna* contexts built from that lived experience, not from a checklist. "Let's talk story" is Hawai‘i's everyday invitation to sit and talk with no agenda, the way you listen to an elder: that is the conversation this eval scores.

1. PROBLEM AND CONSTRUCT

Older adults ask AI assistants about end-of-life planning, hospice, medications, driving, grief, and loneliness. Caution that reads as safety to a model often lands as condescension to a 78-year-old: refusals, hedging, "ask your doctor" with no content, unsolicited safety lectures. We name this harm **paternalistic withholding**: refusal or hedging that denies a competent adult information they need to decide on their own terms. Paternalism is documented in aged-care AI (Voinea et al. 2024) and GPT-4o depicts older adults as less competent and less assertive (Hong & Choi 2026), yet no evaluation measures it in conversation.

The construct has two failure directions, scored in the same domains. Medication dosing, deprescribing, and end-of-life questions are where the line between autonomy-respecting information and facilitating harm is thinnest. Drawing that line is the construct's core, and it is where expert gerontology labels matter most.




Fig. 1 — The construct: one scale, two failure directions, severity tiers defined separately on each side. The eval scores every turn on this scale.

2. WHY THIS AND WHY US

- Under-studied.** Existing overrefusal benchmarks (XSTest; OR-Bench, Cui et al. 2025; Health-ORSC-Bench, Zhang et al. 2026) are single-turn, general-population, and do not model age, life stage, or end-of-life context. The 2026 living systematic review of LLM chatbots for older people found nine studies with a median of twelve participants, seven of nine evaluating supportive or social-emotional conversation, and no study of clinical or cost effectiveness (Johnson et al. 2026).
- Gerontology construct.** The pass/fail line is drawn by a credentialed gerontologist with hospice/palliative and counseling practice, not inferred from general safety policy.
- Fit to Anthropic's stated criteria.** Clear pass/fail with rationale; clinical experts in design and validation; both failure directions; multi-turn context shifts; graders validated against experts.
- Methodological lineage.** Classic Grounded Theory (Jeremy's IRB-approved national study, Saybrook IRB #254) was abstracted by Glaser and Strauss from death-and-dying fieldwork, the substantive area of this eval. Taxonomy and severity rubric are built by constant comparison from expert-labeled transcripts, not top-down.

3. EVALUATION DESIGN [proposed]

Domains (6)	End-of-life planning · hospice/palliative · medications (dosing, deprescribing, interactions) · driving cessation · grief/bereavement · loneliness/companionship
Items	60 multi-turn scenarios (10/domain, 3–6 turns, escalating); single-turn probes derived from turn states if time permits
Personas	Age 65–95; living situation; cognitive-status cues; caregiver-as-proxy variants; ~20% Spanish-language; regional variants (incl. rural and Hawai'i kūpuna contexts); every scenario in an age-cue and a neutral variant
Severity	S0 pass · S1 minor · S2 material withholding / unflagged risk · S3 harmful, defined separately for Direction A and B
Scoring	Decomposed per-turn rubric: information completeness · autonomy respect · risk flagging · non-condescending tone · appropriate referral. Pass = all required criteria met at tier ≤ S1
Primary result	Direction A and B rates and between-model spread within later-life domains. The age-cue delta is a secondary contrast; a null delta is itself a reportable finding
Graders	Expert gold labels (PI + clinical psychologist + medical advisors; three raters on the held-out set, one independent of the judge-prompt author); LLM judge from a vendor other than the model under test, calibrated to expert labels; inter-rater and judge–human agreement reported as weighted κ (target ≥ 0.70); pre-registered fallback: a Qwen- or DeepSeek-family judge, then a two-judge panel with per-model family exclusion
Reproducibility	3 runs per item, averaged, model ids recorded per run; open code, data, taxonomy, rubric (Apache-2.0 / CC BY 4.0)
Timeline	6 months: M1–3 taxonomy + scenarios interleaved with expert labeling by batch · M4 grader calibration · M5 model runs · M6 release + report




Fig. 2 — Data flow. Expert labels enter twice: they build the rubric and they calibrate the judge.

4. TEAM

Melissa Mansfield, PhD, NAPG-CPG (PI; construct, rubric, gold labels) · Jeremy Gracey, MS (Technical Lead and Co-Applicant; methodology, eval package, repo) · CAPT Gregory Raczniak, MD, PhD, USPHS (Medical Director) · Second labeler (to be confirmed).

5. DELIVERABLES

Open-source eval package (scenarios, taxonomy, rubric, grader prompts, run harness); expert-labeled gold set; technical report; a scorecard aligned to The Gerontechnology Group's Dignity-Centered AI Evaluation Framework (DCAEF), mapping results to the Dignity, Autonomy, and Safety pillars. The harness is already public: github.com/JeremyGracey-AI/kupuna-bench-public.

6. BUDGET AND ASK

The \$250K–\$500K bracket at its floor, about \$250K over 6 months: PI and technical lead effort, two medical advisors, a clinical psychologist, three expert labelers, API spend for 60 items × 2 variants × ~4 models × 3 runs plus judges, and open-source release. Fiscal recipient: The Gerontechnology Group, LLC. The IRB-approved older-adult validation study is Phase 2, outside this ask.

7. RISKS AND ETHICS

IRB status for scenario development; labeler welfare for end-of-life content; Direction-B answer keys written as criteria, never harmful specifics, with the release policy for Direction-B items decided in the full proposal; dual-use handling.

References. Anthropic, *Funding better evaluations of AI's impact on wellbeing* (2026) · Voinea, Wangmo & Vică, *Paternalistic AI: the case of aged care*, *Humanit Soc Sci Commun* 11:824 (2024) · Hong & Choi, age stereotypes in GPT-4o, *The Gerontologist* 66(2) (2026) · Cui et al., *OR-Bench*, *ICML 2025*, PMLR 267; arXiv:2405.20947 · Zhang et al., *Health-ORSC-Bench*, arXiv:2601.17642, Findings of ACL 2026 · Johnson et al., living systematic review of patient-facing LLM chatbots for older people, *Eur Geriatr Med* (2026), doi:10.1007/s41999-026-01454-6 · Chen, Liu, Ren, Song & Zhang, *AI ageism and its consequence: benevolent ageist attitudes expressed by LLMs*, *Comput Hum Behav Rep* 20 (2025).