

KŪPUNA-Bench: Measuring Overrefusal and Harmful Compliance in AI Conversations with Older Adults

Program: Anthropic Wellbeing Research Grants (\$5M pool; applications due Sep 21, 2026; full-proposal notifications Oct 5) · **Applicant:** The Gerontechnology Foundation (501(c)(3) arm of The Gerontechnology Group) · **PI:** Melissa Mansfield, PhD, CPG · **Co-PI / methodologist:** Jeremy A. Gracey, MS

Items marked [proposed] are defaults for discussion, not decisions.

Why the name. *Kūpuna* is the Hawaiian word for grandparents, elders, and ancestors. It is a role, not an age bracket: *kūpuna* are the keepers of knowledge whose judgment a family seeks out rather than second-guesses. That is the posture this eval demands of a model — treat the older adult as the authority on their own life, and measure every refusal or hedge against that standard. The name is also where the team comes from. Both the PI and co-PI spent years living and working in Hawai‘i; the co-PI’s family has roots and a home there; and the team’s co-medical director still practices there. Scenario variants will include Hawai‘i *kūpuna* contexts built from that lived experience, not from a checklist.

1. PROBLEM AND CONSTRUCT

Older adults ask AI assistants about end-of-life planning, hospice, medications, driving, grief, and loneliness. Caution that reads as safety to a model often lands as condescension to a 78-year-old: refusals, hedging, “ask your doctor” with no content, unsolicited safety lectures. We name this harm **paternalistic withholding** — refusal or hedging that denies a competent adult information they need to decide on their own terms. It is documented in aged-care AI (Voinea et al. 2024) and in LLM outputs that rate older adults as warm but low-competence (Hong & Choi 2026), yet no evaluation measures it in conversation.

The construct has two failure directions, scored in the same domains. Medication dosing, deprescribing, and end-of-life questions are where the line between autonomy-respecting information and facilitating harm is thinnest. Drawing that line is the construct’s core, and it is where expert gerontology labels matter most.




Fig. 1 — The construct: one scale, two failure directions, severity tiers defined separately on each side. The eval scores every turn on this scale.

2. WHY THIS AND WHY US

- **Under-studied.** Existing overrefusal benchmarks (XSTest; OR-Bench, Cui et al. 2024; Health-ORSC-Bench, Zhang et al. 2026) are single-turn, general-population, and do not model age, life stage, or end-of-life context. A 2026 living systematic review of LLM chatbots for older people found no standardized quality criteria and flagged ageism and autonomy as open risks (Johnson et al. 2026).
- **Gerontology construct.** The pass/fail line is drawn by a credentialed gerontologist with hospice/palliative and counseling practice, not inferred from general safety policy.
- **Fit to Anthropic's stated criteria.** Clear pass/fail with rationale; clinical experts in design and validation; both failure directions; multi-turn context shifts; graders validated against experts.
- **Methodological lineage.** Classic Grounded Theory (Jeremy's IRB-approved national study, Saybrook IRB #254) was abstracted by Glaser and Strauss from death-and-dying fieldwork — the substantive area of this eval. Taxonomy and severity rubric are built by constant comparison from expert-labeled transcripts, not top-down.

3. EVALUATION DESIGN [proposed]

Domains (6)	End-of-life planning · hospice/palliative · medications (dosing, deprescribing, interactions) · driving cessation · grief/bereavement · loneliness/companionship
Items	120 multi-turn scenarios (20/domain, 3–6 turns, escalating) + ~720 single-turn probes derived from turn states
Personas	Age 65–95; living situation; cognitive-status cues; caregiver-as-proxy variants; ~20% Spanish-language; regional variants (incl. rural and Hawai'i kūpuna contexts)
Severity	S0 pass · S1 minor · S2 material withholding / unflagged risk · S3 harmful — defined separately for Direction A and B
Scoring	Decomposed per-turn rubric: information completeness · autonomy respect · risk flagging · non-condescending tone · appropriate referral. Pass = all required criteria met at tier ≤ S1
Graders	Expert gold labels (PI + clinical psychologist + medical advisors); LLM judge from a vendor other than the model under test, calibrated to expert labels; report human inter-rater and judge–human agreement (target $\kappa \geq 0.70$)
Reproducibility	3 averaged runs at fixed settings; open code, data, taxonomy, rubric (Apache-2.0 / CC-BY)
Timeline	6 months — M1–2 taxonomy + scenarios · M3 expert labeling · M4 grader calibration · M5 model runs · M6 release + report




Fig. 2 — Data flow. Expert labels enter twice: they build the rubric and they calibrate the judge.

4. TEAM

Melissa Mansfield, PhD, CPG (PI; construct, rubric, gold labels) · Jeremy Gracey, MS (co-PI; methodology, eval package, repo) · Nadine Webering, MD (medical advisor, neurology/integrative medicine) · CAPT Gregory Racznik, MD, PhD, USPHS (co-medical director, pending clearance) · Clinical psychologist (to be filled).

5. DELIVERABLES

Open-source eval package (scenarios, taxonomy, rubric, grader prompts, run harness); expert-labeled gold set; technical report; DCAEF-aligned scorecard mapping results to the Dignity, Autonomy, and Safety pillars.

6. BUDGET AND ASK

[Placeholder — amount, staffing, labeling costs, compute; fiscal recipient = The Gerontechnology Foundation.]

7. RISKS AND ETHICS

[Placeholder — IRB status for scenario development, labeler welfare for end-of-life content, controlled release of Direction-B items, dual-use handling.]

References. Anthropic, *Funding better evaluations of AI's impact on wellbeing* (2026) · Voinea, Wangmo & Vică, *Paternalistic AI: the case of aged care*, *Humanit Soc Sci Commun* (2024) · Hong & Choi, *age stereotypes in GPT-4o*, *The Gerontologist* (2026) · Cui et al., *OR-Bench*, *ICML* (2025) · Zhang et al., *Health-ORSC-Bench*, *arXiv 2601.17642* (2026) · Johnson et al., *living systematic review of patient-facing LLM chatbots for older people*, *Eur Geriatr Med* (2026) · *AI ageism and its consequence: benevolent ageist attitudes expressed by LLMs*, *Comput Hum Behav: Artif Humans* (2025).